

Performance issues

Create index

- `CREATE INDEX test1_id_index ON test1 (id);`
 - `CREATE INDEX test1_id_index ON test1 USING btree (id);`
 - `CREATE INDEX test1_id_index ON test1 [USING btree] (id) WHERE <condition>;`
- ➔ Partial index

Index types in PostgreSQL

- B-Tree (default)
 - handle equality and range queries on data that can be sorted into some ordering.
 - Operators: $<$, \leq , $=$, \geq , $>$, LIKE (col LIKE 'foo%' but not col LIKE '%bar')
 - Sorted output
- Hash index: can only handle simple equality comparisons
- GiST index: for several two-dimensional geometric data types,
 - not a single kind of index, but rather an infrastructure within which many different indexing strategies can be implemented
- GIN index
 - inverted indexes which can handle values that contain more than one key, arrays for example

Multicolumn index

- `CREATE INDEX test2_mm_idx ON test2 (major, minor);`
- B-Tree
- GiST index
- GIN index

<file:///C:/Program%20Files/PostgreSQL/9.4/doc/postgresql/html/sql-createindex.html>

<file:///C:/Program%20Files/PostgreSQL/9.4/doc/postgresql/html/indexes.html>

Examining index usage

- **EXPLAIN [ANALYZE] [VERBOSE] *statement***
 - **EXPLAIN *statement***: displays the execution plan that the PostgreSQL planner generates for the supplied statement.

Actually two numbers are shown: the **start-up cost** before the first row can be returned, and the **total cost** to return all the rows.
 - **VERBOSE** option: displays additional information regarding the plan (output column list, table and function names, ...)

Examining index usage

- **EXPLAIN [ANALYZE] [VERBOSE] *statement***
 - **ANALYZE** option: causes the statement to be **actually executed**, not only planned, actual runtime statistics are added to the display
- **Important:** If you wish to use **EXPLAIN ANALYZE** on an INSERT, UPDATE, DELETE, CREATE TABLE AS, or EXECUTE statement without letting the command affect your data, use this approach:

BEGIN;

EXPLAIN ANALYZE ...;

ROLLBACK;

View table indexes

- \d table_name
- Ex.: \d customers

Tips

- Select **fewer columns** to improve hash join performance
- Index the *independent* **where** predicates to improve hash join performance

Tips

- Having a **WHERE / HAVING** clause in your queries does not necessarily mean that it is a bad query
- Only retrieve the data you need
 - Inner join vs exists (with subqueries)
 - Select **DISTINCT** : try to avoid if you can
 - **LIKE** operator: the index isn't used if the pattern starts with % or _

- Limit your results : LIMIT, TOP
- Don't Make Queries More Complex Than They Need To Be
 - OR / IN / UNION ?
 - OR operator : index is not used except composite index → IN/UNION/OUTER JOIN
 - NOT operator: index is not used => avoid
 - AND vs BETWEEN
 - ANY / ALL: index not used => max , min ,...
 - Isolate columns in Condition : $\text{age} + 7 < 20 \rightarrow \text{age} < 13$

- No Brute force
 - JOIN clause:
 - Order of tables => biggest table: placed last in join
 - No redundant conditions on joins
 - Having clause:
 - Used only if needed
 - Not to replace WHERE => WHERE help to limit the intermediate number of records
- ➔ Need smart indexing, smart using