



TRƯỜNG ĐẠI HỌC BÁCH KHOA HÀ NỘI
HANOI UNIVERSITY OF SCIENCE AND TECHNOLOGY

Database

Lesson 12. Transaction management

Ba Lam Do

Learning Map

Sequence	Title
1	Introduction to Databases
2	Relational Databases
3	Relational Algebra
4	Structured Query Language – Part 1
5	Structured Query Language – Part 2
6	Constraints and Triggers
7	Entity Relationship Model
8	Functional Dependency
9	Normalization
10	Storage - Indexing
11	Query Processing
12	Transaction Management – Part 1
13	Transaction Management – Part 2

Outline

- Transaction
 - Definition
 - ACID properties
 - Transaction states
 - Transaction management interfaces
- Concurrency control
 - Objective
 - Scheduling
 - Serializability
 - Locks
 - Locking protocol
 - 2 Phase Locking (2PL)

Objectives

- Upon completion of this lesson, students will be able to:
 - Understanding transaction and ACID properties
 - Understanding locking and 2 phase locking algorithm

1. Transaction

- Definition
- ACID properties
- Transaction states
- Transaction management interfaces

1.1. What is transaction

- A transaction is a unit of program execution that accesses and possibly updates various data items.
- Example: transfer \$50 from account A to account B
 - 1. $R(A)$
 - 2. $A \leftarrow A - 50$
 - 3. $W(A)$
 - 4. $R(B)$
 - 5. $B \leftarrow B + 50$
 - 6. $W(B)$
- Two main issues:
 - 1. concurrent execution of multiple transactions
 - 2. failures of various kind (e.g., hardware failure, system crash)

1.2. ACID Properties

- Database system must guarantee ACID for transactions:
 - **Atomicity**: either all operations of the transaction are executed or none
 - **Consistency**: execution of a transaction in isolation preserves the consistency of the database
 - **Isolation**: although multiple transactions may execute concurrently, each transaction must be unaware of the other concurrent transactions.
 - **Durability**: After a transaction completes successfully, changes to the database persist even in case of system failure.

1.2.1. Atomicity

- Example: transfer \$50 from account A to account B
 1. $R(A)$
 2. $A \leftarrow A - 50$
 3. $W(A)$
 4. $R(B)$
 5. $B \leftarrow B + 50$
 6. $W(B)$
- What if failure (hardware or software) after step 3?
 - money is lost
 - database is inconsistent
- **Atomicity:**
 - either all operations or none
 - updates of partially executed transactions not reflected in database

1.2.2. Consistency

- Example: transfer \$50 from account A to account B
 1. $R(A)$
 2. $A \leftarrow A - 50$
 3. $W(A)$
 4. $R(B)$
 5. $B \leftarrow B + 50$
 6. $W(B)$
- **Consistency in example:** sum $A + B$ must be unchanged
- **Consistency in general:**
 - explicit integrity constraints (e.g., foreign key)
 - implicit integrity constraints (e.g., sum of all account balances of a bank branch must be equal to branch balance)
- **Transaction:**
 - must see consistent database
 - during transaction inconsistent state allowed
 - after completion database must be consistent again

1.2.3. Isolation

- Example: transfer \$50 from account A to account B
 - 1. $R(A)$
 - 2. $A \leftarrow A - 50$
 - 3. $W(A)$
 - 4. $R(B)$
 - 5. $B \leftarrow B + 50$
 - 6. $W(B)$
- **Imagine second transaction $T2$:**
 - $T2 : R(A), R(B), \text{print}(A + B)$
 - $T2$ is executed between steps 3 and 4
 - $T2$ sees an inconsistent database and gives wrong result


1.2.3. Isolation

- **Trivial isolation:** run transactions serially
- Isolation for concurrent transactions:
 - For every pair of transactions T_i and T_j ,
 - it appears to T_i as if either T_j finished execution before T_i started
 - or T_j started execution after T_i finished.

1.2.4. Durability

- When a transaction is done it commits.
- **Example:** transaction commits too early
 - transaction writes A, then commits
 - A is written to the disk buffer
 - then system crashes
 - value of A is lost
- **Durability:** After a transaction has committed, the changes to the database persist even in case of system failure.
- **Commit only after all changes are permanent:**
 - either written to log file or directly to database
 - database must recover in case of a crash

1.3. Transaction states



1.4. Transaction management interfaces

- Begin Trans
- Commit ()
- Abort()
- Savepoint Save()
- Rollback (savepoint)
(savepoint = 0 ==> Abort)

Schedule

- A list of operations conducted by concurrent transactions
- Serial schedule: transactions are executed non-interleaved, i.e., a serial schedule is one in which no transaction starts until a running transaction has ended
- Non-serial schedule: operations of multiple transactions are interleaved

Schedule

T ₁	T ₂
R(A)	
W(A)	
R(B)	
	W(B)
	R(A)
	R(B)

Serial schedule

T ₁	T ₂
R(A)	
W(A)	
	W(A)
	R(A)
commit	
	commit

Non-serial schedule

Problem 1 – Lost updated data

- $A = 100 + 100 - 10 = 190$
- Updated data is overwritten by other transaction

T₁	T₂	Account A
	Begin Transaction	100
Begin Transaction	Read(A);	100
Read(A);	$A = A + 100 ;$	100
$A := A - 10;$	Write(A);	200
Write(A);	Commit	90
Commit		90

Problem 2 – Uncommitted transaction

- $A = 100 - 10 = 90$
- A transaction allows other transactions to see temporary results before they are committed

T₃	T₄	TK A
	Begin Transaction	100
	Read(A);	100
	$A = A + 100 ;$	100
Begin Transaction	Write(A);	200
Read(A);	:	200
$A := A - 10;$	Rollback	100
Write(A);		190
Commit		190

Problem 3 – Dirty read/Unrepeatable read

- Read-only transaction reads data which is being updated by other transaction

T₅	T₆	X	Y	Z	Sum
	Begin Transaction	100	50	25	
Begin Transaction	Sum = 0	100	50	25	0
Read(X);	Read(X);	100	50	25	0
X = X - 10 ;	Sum = Sum + X	100	50	25	100
Write (X);	Read(Y);	90	50	25	100
Read (Z)	Sum = Sum + Y	90	50	25	150
Z = Z + 10		90	50	25	150
Write(Z);		90	50	35	150
Commit	Read(Z)	90	50	35	150
	Sum = Sum + Z	90	50	35	185
	Commit	90	50	35	185

2. Concurrency control

- Objective
- Scheduling
- Serializability
- Locks
- Locking protocol
- 2 Phase Locking (2PL)

2.1. Objective

- Objective:
 - ensures that database transactions are performed concurrently without violating the data integrity
 - guarantees that no effect of committed transactions is lost, and no effect of aborted (rolled back) transactions remains in the related database.
- Example

<pre>T0: read(A); A := A - 50; write(A); read(B); B := B + 50; write(B);</pre>	<pre>T1: read(A); temp := A * 0.1; A := A - temp; write(A); read(B); B := B + temp; write(B);</pre>
--	---

2.1. Scheduling

T ₀	T ₁
read(A) A := A - 50 write(A) read(B) B := B + 50 write(B)	read(A) temp := A * 0.1 A := A - temp write(A) read(B) B := B + temp write(B)

(1)

T ₀	T ₁
read(A) A := A - 50 write(A) read(B) B := B + 50 write(B)	read(A) temp := A * 0.1 A := A - temp write(A) read(B) B := B + temp write(B)

(2)

T ₀	T ₁
read(A) A := A - 50 write(A) read(B) B := B + 50 write(B)	read(A) temp := A * 0.1 A := A - temp write(A) read(B) B := B + temp write(B)

(3)

2.2. Serializability

- A schedule of a set of transactions is a linear ordering of their actions
 - e.g. for the simultaneous deposits example:
 $R1(X) R2(X) W1(X) W2(X)$
- A serial schedule is one in which all the steps of each transaction occur consecutively
- A serializable schedule is one which is equivalent to some serial schedule

2.3. Locks

- A lock is a mechanism to control concurrency on a data item.
 - one way of enforcing concurrency control policies
- Two types of locks on a data item A :
 - **exclusive** – $xL(A)$: prevents other transactions from accessing the data.
 - **shared** – $sL(A)$: data can only be read.
- Lock request are made to **concurrency control manager**.
- Transaction is blocked until lock is granted.
- **Unlock** A – $uL(A)$: release the lock on a data item A

2.3.1. Lock Compatibility

- Lock compatibility matrix:

$T_1 \downarrow T_2 \rightarrow$	shared	exclusive
shared	true	false
exclusive	false	false

- **T_1 holds shared lock on A:**
 - shared lock is granted to T_2
 - exclusive lock is not granted to T_2
- **T_2 holds exclusive lock on A:**
 - shared lock is not granted to T_2
 - exclusive lock is not granted to T_2
- **Shared locks can be shared by any number of transactions.**

2.4. Locking protocol

- Example transaction T_2 with locking:
 1. $sL(A), R(A), uL(A)$
 2. $sL(B), R(B), uL(B)$
 3. $print(A + B)$
- **T_2 uses locking, but is not serializable**
 - A and/or B could be updated between steps 1 and 2
 - printed sum may be wrong
- **Locking protocol:**
 - set of rules followed by all transactions while requesting/releasing locks
 - locking protocol restricts the set of possible schedules

Pitfalls of Locking Protocols – Deadlock


- **Example:** two concurrent money transfers
 - $T1: R(A), A \leftarrow A + 10, R(B), B \leftarrow B - 10, W(A), W(B)$
 - $T2: R(B), B \leftarrow B + 50, R(A), A \leftarrow A - 50, W(A), W(B)$
 - possible concurrent scenario with locks:
 - $T1.xL(A), T1.R(A), T2.xL(B), T2.R(B), T2.xL(A), T1.xL(B), \dots$
 - $T1$ and $T2$ block each other – no progress possible

2.4.2. Pitfalls of locking protocols – Starvation

- **Starvation:** transaction continues to wait for lock
- **Examples:**
 - the same transaction is repeatedly rolled back due to deadlocks
 - a transaction continues to wait for an exclusive lock on an item while a sequence of other transactions are granted shared locks
- Well-designed concurrency manager avoids starvation.

2.5. 2 Phase Locking (2PL)

- Protocol that guarantees serializability.
 - Phase 1
 - locks are acquired and no locks are released
 - Phase 2
 - locks are released and no locks are acquired



Two-Phase Locking – Example

- Example: two concurrent money transfers
 - $T1: R(A), A \leftarrow A + 10, R(B), B \leftarrow B - 10, W(A), W(B)$
 - $T2: R(A), A \leftarrow A - 50, R(B), B \leftarrow B + 50, W(A), W(B)$
- Possible two-phase locking schedule:
 - 1. $T1 : xL(A), xL(B), R(A), R(B), W(A \leftarrow A + 10), uL(A)$
 - 2. $T2 : xL(A), R(A), xL(B)$ (*wait*)
 - 3. $T1 : W(B \leftarrow B - 10), uL(B)$
 - 4. $T2 : R(B), W(A \leftarrow A - 50), W(B \leftarrow B + 50), uL(A), uL(B)$
- Equivalent serial schedule: $T1, T2$

Solve Problem 1 - Lost updated data

Time	T ₁	T ₂	TK A
t ₁		Begin Transaction	100
t ₂	Begin Transaction	Write_lock(A)	100
t ₃	Write_lock(A)	Read(A);	100
t ₄	WAIT	A = A + 100 ;	100
t ₅	WAIT	Write(A);	200
t ₆	WAIT	Commit/ unlock(A)	200
t ₇	Read(A);		200
t ₈	A:=A-10;		190
t ₉	Write(A);		190
t ₁₀	Commit/ Unlock(A)		190

write lock: the relevant transaction can read and write data

read lock: the relevant transaction can read but can not write data

Solve Problem 2 – Uncommitted transaction

Time	T ₃	T ₄	TK A
t ₁		Begin Transaction	100
t ₂		Write_lock (A)	100
t ₃	Begin Transaction	Read(A);	100
t ₄	Write_lock (A)	A = A + 100 ;	100
t ₅	WAIT	Write(A);	200
t ₆	WAIT	Rollback/ unlock(A)	100
t ₇	Read(A);		100
t ₈	A:=A-10;		90
t ₉	Write(A);		90
t ₁₀	Commit/ Unlock(A)		90

write lock: the relevant transaction can read and write data

read lock: the relevant transaction can read but can not write data

Solve Problem 3 – Dirty read

t	T₅	T₆	X	Y	Z	Sum
t ₁		Begin Transaction	100	50	25	
t ₂	Begin Transaction	Sum = 0	100	50	25	0
t ₃	Write lock (X)	Read lock (X)	100	50	25	0
t ₄	Read(X);	WAIT	100	50	25	0
t ₅	X = X - 10 ;	WAIT	100	50	25	0
t ₆	Write (X);	WAIT	90	50	25	0
t ₇	Write lock (Z)	WAIT	90	50	25	0
t ₈	Read (Z)	WAIT	90	50	25	0
t ₉	Z = Z + 10	WAIT	90	50	25	0
t ₁₀	Write(Z);	WAIT	90	50	35	0
t ₁₁	Commit/ Unlock(X,Z)	WAIT	90	50	35	0
t ₁₂		Read(X);	90	50	35	0
t ₁₃		Sum = Sum + X	90	50	35	90
t ₁₄		Read lock(Y)	90	50	35	90
t ₁₅		Read(Y);	90	50	35	90
t ₁₆		Sum = Sum + Y	90	50	35	140
t ₁₇		Read lock(Z)	90	50	35	140
t ₁₈		Read(Z)	90	50	35	140
t ₁₉		Sum = Sum + Z	90	50	35	175
t ₂₀		Commit/ Unlock(X,Y,Z)	90	50	35	175

Summary

- Transaction and ACID properties
- Concurrency control based on locking mechanisms
- 2 phase locking (2PL)

Next lesson: Transaction management

- Hector Garcia-Molina, Jeffrey D. Ullman, Jennifer Widom. Database Systems: The Complete Book. Pearson Prentice Hall. the 2nd edition. 2008: Chapter 7
- Nguyen Kim Anh, Nguyên lý các hệ cơ sở dữ liệu, NXB Giáo dục. 2004: Chương 7